Master AI Career Document: Tyler Stahl

Core Repository for Professional Mission Statements, Applied AI Résumé Elements, and Targeted Cover Letters (Ramp / Google Cloud)

SECTION 1: MASTER MISSION STATEMENTS

Use these variations to quickly paste your identity into LinkedIn connection notes, emails, or portfolio introductions depending on the space allowed.

Short Version (LinkedIn Connection Notes / Max 300 Characters)

Full-stack engineer and applied AI builder specializing in multi-agent orchestration. Combining 7+ years scaling web platforms for millions with an MA in Counseling to bridge the human-to-AI gap, creating advanced autonomous systems into trusted, intuitive tools stakeholders confidently adopt.

Medium Version (Email Blurb / Paragraph)

I am a full-stack software engineer and applied AI developer who acts as the connective bridge between complex autonomous architecture and the human stakeholders who use it. Over a seven-year engineering career scaling data-dense portals for millions of users, I have learned that the ultimate bottleneck for advanced AI is not computational power—it is human trust. Leveraging an interdisciplinary background that includes a Master's degree in Counseling, I build the Human-in-the-Loop workflows, stable specs, and technical change management strategies that turn volatile AI systems into reliable, user-centric enterprise tools.

Long Version (Resume Summary / Executive Bio)

Full-Stack Software Engineer and Applied AI Developer specializing in bridging advanced multi-agent workflows with intuitive, human-centric production environments. Bringing over seven years of professional engineering experience scaling data-dense web infrastructure for millions of annual users, my focus is on ensuring complex autonomous systems deliver practical utility and inspire user confidence. By combining hands-on technical execution—including structured extraction pipelines, hybrid LLM inference routing, and multi-agent orchestrations—with an academic background in behavioral science and anthropology, I design resilient, high-velocity products that prioritize organizational trust, transparent change management, and seamless human-agent collaboration.

SECTION 2: THE CORE APPLIED AI RÉSUMÉ

Optimized with quantified metrics, marketing partnerships, and your humanities advantage.

Contact Information

Tyler Stahl Glen Rock, NJ 07452 | (206) 402-1346 | tylerstahl@gmail.com | linkedin.com/in/tyler-j-stahl

Core Competencies

- **Applied AI & Infrastructure:** Multi-Agent Orchestration, LLM Inference (Cloud/Local), Structured Data Extraction, Prompt Engineering, Autonomous Workflows.
- **Full-Stack Engineering:** React, TypeScript, Python, Laravel, JavaScript, HTML5/CSS, TailwindCSS.
- **Backend & Data:** GraphQL, REST APIs, SQLite, PostgreSQL, Data Pipelines (ETL).
- **Engineering Processes:** CI/CD Pipelines, Docker, Automated Testing (Cypress, Playwright, Jest), Git, Agile Delivery.

Professional Experience

Independent R&D | NJ

Applied AI Developer | February 2026 – Present
Building advanced AI infrastructure, recently shipping a sophisticated multi-agent orchestration ecosystem.

- **Multi-Agent Orchestration:** Built and deployed a secure, distributed intelligence pilot utilizing the OpenClaw framework to manage a workforce of specialized autonomous agents (Orchestrator, Data Analyst, UI Developer).
- **LLM Inference & Infrastructure:** Engineered a hybrid inference strategy, routing complex reasoning tasks to Google Gemini 3.1 Pro and local Gemma 27b, optimizing operational token costs by 40% by routing high-velocity parsing tasks to Gemini 2.5 Flash and Gemma 4b.
- **Structured Extraction & Data Pipelines:** Developed autonomous Python ETL pipelines to scrape, clean, and structure complex market data, utilizing SQLite for durable state management and Google Workspace CLI for automated reporting.
- **Internal Tooling & Security:** Hardened the agentic environment by building Human-in-the-Loop (HITL) terminal approvals via Antigravity IDE, implementing strict subagent gateway whitelisting, and managing secrets via 1Password.

National Alliance on Mental Illness (NAMI)

Front-End Developer | December 2024 – January 2026
Contributed to core infrastructure and internal tooling to scale front-end development across the organization.

- **Platform Infrastructure:** Built and implemented a centralized Design Component Library, helping enforce architecture standards and streamline UI deployment for the core business application.
- **Engineering Process Improvement:** Developed automated End-to-End (E2E) testing pipelines utilizing Cypress, reducing manual QA cycles and improving deployment confidence.
- **Quality & Compliance:** Executed comprehensive code quality and accessibility audits, collaborating on cross-functional technical change initiatives that increased the aggregate compliance score by 30%.

Charity Navigator | NJ

Software Engineer | August 2023 – September 2024
Contributed to the delivery of a massive-scale data platform in a high-traffic environment.

- **Full-Stack Production Scale:** Engineered features for a React/TypeScript portal serving over 50,000 clients, optimizing front-end architecture to sustain high-traffic events of upwards of 700,000 concurrent visitors.
- **Backend Integration:** Assisted in the complex transition of legacy data flows into optimized GraphQL API endpoints, streamlining backend-to-frontend communication.

Cross-Functional Collaboration: Worked directly with product and data teams to ship tools that drove operational excellence and streamlined data review processes.

Charity Navigator | NJ

Front-End Web Developer | March 2019 – August 2023

- **Strategic Marketing Partnership:** Partnered directly with the Chief Marketing Officer and the marketing team to execute cross-functional initiatives, ranging from large-scale strategic projects like the organization's 20th-anniversary rebrand to the rapid deployment of time-sensitive website features aligned with strict marketing schedules.
- **High-Impact Delivery:** Played a key role in the technical execution of a massive platform rebuild, collaborating with external design agencies to transition mockups into a production-ready system for over 11 million annual users.
- **Business Outcomes:** Delivered full-stack technology initiatives that contributed to securing a $5M grant and increased processing efficiency by 43%.

Volant Web Design | NJ

Founder & Web Developer | April 2017 – March 2019 & September 2024 – Present

- Shipped complete full-stack software solutions (PHP, JS, HTML/CSS) for diverse B2B clients, managing the entire lifecycle from requirements to production deployment and SEO optimization.

Education

- **Front-End Nanodegree**, Udacity (Awarded full scholarship from Google's Grow with Google program)
- **Master of Arts (MA), Counseling**, Seattle University
    - *(Identity Hook: Leveraged to architect Human-in-the-Loop workflows, ensuring organizational trust and responsible AI deployment).*
- **Bachelor of Arts (BA), Anthropology**, Dartmouth College

SECTION 3: MASTER COVER LETTERS

Document A: Google Cloud - Forward Deployed Engineer (FDE)

Dear Google Cloud Hiring Team,

Google Cloud CEO Thomas Kurian's announcement regarding a massive initiative to deploy Forward Deployed Engineers (FDEs) to scale customer AI and agentic development underlines a critical industry truth: the frontier of AI is no longer about model size alone, but about the practical orchestration layer that integrates it into business workflows. I am applying for the Forward Deployed Engineer position because I am a fundamentally practical builder who specializes in architecting the production scaffolding that takes advanced LLM capabilities out of isolation and embeds them safely into scalable enterprise solutions.

My recent independent R&D has focused heavily on the exact multi-agent orchestration and Python mastery Google is prioritizing for this scaling effort. Utilizing the OpenClaw framework, I built and deployed a secure, distributed intelligence pilot managing a coordinated workforce of specialized autonomous sub-agents. To address the enterprise challenges of cost and latency optimization, I engineered a tri-tier hybrid inference routing strategy. This architecture routes high-complexity reasoning tasks to Google Gemini 3.1 Pro, while dynamically offloading high-velocity utility parsing tasks to Gemini 2.5 Flash and local Gemma instances. Furthermore, I developed autonomous Python ETL pipelines with SQLite state management to solve for long-running data persistence, and hardened the environment with Human-in-the-Loop (HITL) terminal approval gateways via Antigravity IDE.

Enterprise AI deployment requires more than standalone code; it demands an engineer who understands production-grade scale and high-stakes cross-functional communication. Over my 6-year software engineering career, I have proven my ability to scale infrastructure for millions of users. At Charity Navigator, I optimized a frontend and data architecture that sustained high-traffic spikes of up to 700,000 concurrent visitors for a platform serving 11 million annual users. Crucially, I partnered directly with the Chief Marketing Officer and cross-functional data teams to execute high-velocity tactical deployments and major organizational rebrands. I understand how to sit with non-technical stakeholders, translate complex systemic limitations, and quickly ship minimum viable products (MVPs) that drive operational excellence. This maps directly into the client-facing, high-exposure nature of Google Cloud's Go-To-Market engineering organization.

What sets me apart from traditional software engineers is my interdisciplinary background as a Dartmouth College Anthropology graduate holding a Master's degree in Counseling. I view my humanities background as a technical force multiplier for agentic deployment. I do not just write script; I design the "human-to-agent" workflows, audit trails, and technical change management systems that bridge the trust gap for enterprise stakeholders moving toward automated governance. I am eager to operate as the high-agency "connective tissue" that Google Cloud needs to turn enterprise AI friction points into seamless digital labor models. Given the velocity with which Google's FDE organization is moving to fulfill this customer-facing surge, I am prepared to enter your accelerated interview loop immediately.

Thank you for your time and consideration.

Best regards,
Tyler Stahl

***

Document B: Ramp - Applied AI Engineer

Dear Hiring Team,

Ramp’s commitment to high agency, urgency, and shipping infrastructure that manages billions in business spend requires engineers who are fundamentally practical builders. I am applying for the Applied AI Engineer role because I thrive at the exact intersection Ramp operates in: taking advanced LLM capabilities out of the research phase and integrating them into highly scalable, production-ready web platforms.

With over 6 years of professional experience as a Full-Stack Software Engineer, I know how to ship code to production in high-stakes, data-dense environments. At Charity Navigator, I contributed to the delivery of a React and TypeScript data portal that scaled to support over 50,000 clients and sustained traffic spikes of 700,000 concurrent visitors. I assisted in complex data integrations via GraphQL and collaborated directly with product teams to drive operational excellence. Later, at NAMI, I focused on developer productivity by building automated Cypress testing pipelines and centralized UI component libraries to scale our engineering efforts.

Recently, I have focused my independent R&D deeply on applying LLMs to real-world problems. I architected an advanced multi-agent ecosystem using the OpenClaw framework, designing a system that utilizes a hybrid inference strategy—intelligently routing deep reasoning tasks to Google Gemini Pro and local Gemma 27b models, while optimizing latency by sending lightweight parsing tasks to Gemma 4b. To solve for structured extraction and state management, I built autonomous Python ETL pipelines backed by SQLite to ensure durable memory across agent workflows. I don't just prototype; I build the practical scaffolding (like secure API integrations and HITL terminal approvals) that makes LLMs actually useful.

What sets me apart is my non-traditional background. Before engineering, I earned a Master's degree in Counseling. This foundation makes me deeply collaborative and communicative. I don't just write code in a silo; I work alongside stakeholders to ensure the tools we build actually solve the underlying human and business problems. I combine the rigorous production experience of a full-stack engineer, a strong collaborative background, and a deep obsession for applied AI. I am eager to bring my high-agency mindset to Ramp to help build tools that make finance teams infinitely more productive.

Thank you for your time and consideration. I would welcome the opportunity to discuss how my background aligns with your engineering goals.

Best regards,
Tyler Stahl

–

Applicaiton/Interview questions

How would you describe your unique value as someone who bridges technical AI delivery with human-centered design?
I combine a foundation in humanities with engineering experience. Full-stack engineer and applied AI builder specializing in multi-agent orchestration. Combining 7+ years scaling web platforms for millions with an MA in Counseling to bridge the human-to-AI gap, creating advanced autonomous systems into trusted, intuitive tools stakeholders confidently adopt. I also learned brand standards and customer experience from the best, Apple, while working in their B2B team at the 5th Avenue Apple Store.

Walk me through your multi-agent orchestration pilot: What business problem did it solve and what were the measurable outcomes?

The problem was synthesizing available but fragmented market data into long-term market context and decision making. I set up one agent to pull in market data such as cost of raw materials and cost of finished products, another to run calculations on profitability based on long term trends, another to build a web dashboard of data, and an orchestrator above all to monitor and coordinate. Result was quicker decision making and around 15% increased profit.

Beyond Gemini, which GCP services have you used hands-on — Vertex AI, BigQuery, GKE, or others?

I migrated my projects from an initial POC in Google AI Studio into a bespoke codebase in ANtigravity powered by Vertex AI.

I use Firestore for storing and protecting customer data.

I use IAM for coordinating my builder account access and support agent access.

Tell me about a time you had to manage competing stakeholder expectations on a technical project.

WOrking with an internal client on a pilot meant to scale across multiple departments. My goal was to create modular, reusable components to create custom solutions that could scale and iterate quickly. This required managing customer expectations and product scope so we could deliver on time and not get bogged down in one-off client needs. But I also wanted to 'surprise and delight' the customer, for the delivered product to feel empowering. This took careful framingn of the project as 'ground floor', customer getting first access to a new way of doing thing, and also helping pave the way for their peers to get a robust product in the future. Nobody wants to be told 'no', so I worked on reframing as 'let's look at that and see where it fits'.

See recommendation